

The vibes that I summoned...

Thoughts on working
with AI agents.

Chris Pahl • 2026



The Spectrum (on Reddit)



Full manual

every keystroke yours

Full vibecode


just make it work

← more control · more understanding | more productivity · more delegation →

General tendency:

Using GenAI is a tradeoff between control and productivity.

My take: Quality vs Productivity



Quality peaks in the middle. Productivity follows it down. Confidence diverges from both -> Dunning-Kruger zone!

But there's already data:

-19%

slower with AI

experienced devs · own mature
repos · predicted +24%

METR · RCT · 2025

~40%

insecure programs

Copilot output across 89
security-relevant scenarios

NYU CCS · 2021

8×

more duplicated code

block-level clones up ·
refactoring 24% → 10%

GitClear · 211M LoC · 2024

-7.2%

delivery stability

drop with AI adoption · 39%
distrust AI code

DORA / Google · 2024

29.5%

Python with CWEs

AI-generated code in real
GitHub repos · 43 CWE
categories

Fu et al. · ACM TOSEM · 2025



trust = less scrutiny

more trust in GenAI ⇒ less
critical thinking applied

*Lee et al. · Microsoft + CMU ·
CHI 2025*

17%

Good use cases

rubber-ducking ideation

explain unfamiliar code test generation

pair programming boilerplate refactoring assist

learning accelerator documentation drafts naming things

summarisation regex / SQL crafting edge-case brainstorming

translation code review companion mock data / fixtures

error decoding search engine log analysis proofreading

automation

Risks for all

ecological cost **Convincing hallucinations**
schools licensing **Atrophy** concentration
misinformation at scale **Prompt injection & MCP**
content degradation **No juniors hired** hardware crisis
operator bias recursive collapse economic transfer
data privacy cyberattack automation communities thinning out
erosion of trust

25%

Risks for us

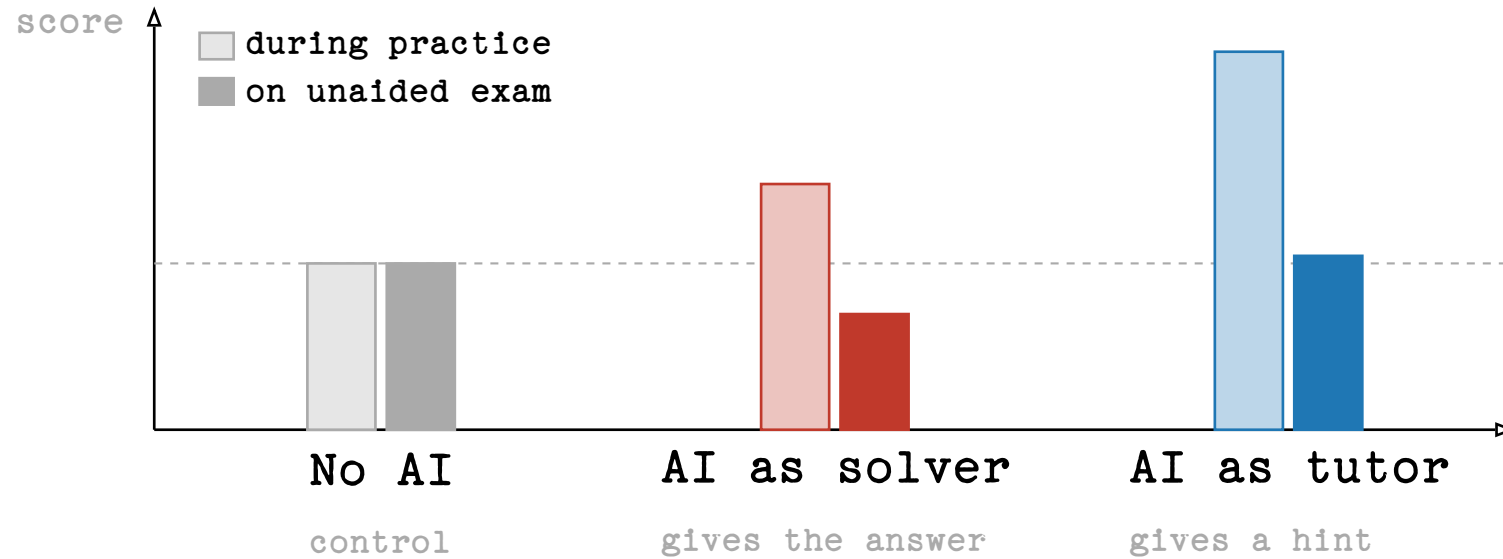
ecological cost **Convincing hallucinations**
schools licensing **Atrophy** concentration
misinformation at scale **Prompt injection & MCP**
content degradation **No juniors hired** hardware crisis
operator bias recursive collapse economic transfer
data privacy cyberattack automation communities thinning out
erosion of trust

Exhibit A: schools.

*Maybe it's not so much about the tool —
but about how we use it?*

33%

Same tool, opposite outcomes



Bastani et al., *Generative AI Can Harm Learning*, SSRN 2024

Are *you* team tutor or team
solver?

42%

Keeping up with Claude

Our team's Claude Code account · April 2026

98.7%

suggestions accepted as-is

1.3% rejected. Is that a review process?

27,757

lines accepted last month

≈ 925 LoC/day · careful review ≈ 100–200 LoC/
hour

Reviewing 27,757 lines carefully ≈ 138 h.

Human cognitive biases

Automation bias

We accept machine suggestions more readily than human ones.

click accept • review later • maybe

Anchoring

The first suggestion shapes the solution space – even when it's wrong.

the model frames the problem before you do

Confirmation bias

We hear what we already believe.

the prompt contains the answer we want

Authority bias

Eloquent + fast + confident reads as expert.

fluency ≠ correctness

Dunning–Kruger

We mistake competence we observe for competence we have.

“this is what I would have written”

Illusion of explanatory depth

We think we understand – until asked to explain.

“yes, I read the diff” → bug

Statistical model biases

Sycophancy

It's easy to talk the model out of a correct answer.

echo chambers – but for code

Historical / representation bias

Trained on the web – heavily modern, English, Western.

defaults track what's common, not what's right

Attention bias

First and last things dominate; the middle evaporates.

rules buried in long context get ignored

Omission bias

Rare and novel answers get suppressed by common ones.

the model is bad at being weird

Hen & Egg

1. If you don't know what good software looks like –
how do you write the right prompt?
2. If you can't understand what the model just generated –
how do you verify it?
3. If you can't code (anymore) –
how can you understand the diffs?

Humans are still very much required.

How do *you* keep up with Claude?


Five habits that helped me

- 1 Sketch first, generate then.
- 2 Split the work - keep some of it manual.
- 3 Have a real verification strategy.
- 4 Don't make yourself replaceable.
- 5 Treat the AI like a colleague.

1. Sketch first, generate then

You set the anchor, you stay in the loop.

- Sketch the SQL query yourself,
Then ask the model to review and optimise.
- Write the function signature and the docstring.
Then let the model fill the body.
- Put TODO comments in your code
Then let Claude work on them.



2. Split the work, do some manual

The parts you write are the parts you understand.

- Claude does not perform well, if given too big scopes.

Split tasks up and prompt them individually.

- Claude can help you split up work.


It just tends to work on the task right away.

- It is easy to lose understanding without noticing.

Keep doing important things manually!

- Large diffs make it easy to get lost.

Commit often so diffs stay reviewable.



3. Verify strategy before code

There are no shortcuts to quality.

Before you accept a generated change

name the signal that would tell you it's wrong.

Comparison against a
reference
implementation

Property-based
tests, fuzzing, ...

Type checkers,
linters, static
analysers

/review and manual
review by colleagues

Replay against real
production data

Work actively on the
code

Example: Noise library for Dart – I don't speak Dart.

4. Don't make yourself replaceable.

You only get replaced if you make yourself replaceable.

- Code was the source of truth before, that's changing.
Now it is context given via design documents.
- Spend the time Claude saves you on the parts it's bad at.
Namely: decisions, judgement, context, design.
- Reading code is much more important now than writing code.
Keep your tools sharp. Be able to work without Claude.

Design · judgement · context · decisions	<i>you</i>
Reading · understanding · review	<i>you</i>
Refactors · idioms · routine logic	<i>contested</i>
Boilerplate · scaffolding · glue	<i>AI</i>
Syntax · typing · trivial lookups	<i>AI</i>



5. Treat Claude like a colleague

Talk to it like your seat neighbor.

- Explain the tasks at hand like you would to a junior colleague.

A very eager, junior colleague with seemingly infinite capacity.

- Push back. Ask for alternatives. Let it explain. Disagree.

If you'd reject a colleague's PR for that reasoning, reject the model's.

 you commented

Why a map here and not a for-loop?

 claude commented

Map is more idiomatic – happy to switch if perf matters.

 you commented

Show me the benchmark first.

Bonus: Programming as Theory Building

Peter Naur, 1985

- The code is a by-product. What you're really building is the theory – your team's shared mental model of the problem.
- Documentation captures the artefact, not the theory. The bugs, the dead ends, the “why not this?” decisions live in people's heads.
- This is why handovers fail. The theory doesn't transfer cleanly.




Peter Naur, 1928–2016

Photo: Wikimedia Commons

Bonus: What that means for Claude

- The model has no theory – it sees the artefact and fills the rest in.
Confidently. It won't tell you which is which.
- Bad prompt: “Write tests for this file”
It freezes today's behaviour, bugs included. Again, no mention.
- Your job is still to build the theory.
The AI helps you write the code that follows from it.



Questions or Disagreements?

Homework:

PROGRAMMING AS THEORY BUILDING.

